



## Practical no – 1

On the fruit dataset, compare the performance of Logistic Regression, SVM, KNN on the basis of their accuracy.

```
sns.countplot(df['fruit_name'],label='Count') # count plot  
plt.show()
```



```
df.drop('fruit_label',axis=1).plot(kind='box', subplots=True, layout=(2,2),  
sharex=False, sharey=False, figsize=(6,6), title='Box Plot for each input variable')  
plt.savefig('fruits_box')  
plt.show()
```



```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
feature_names = ['mass', 'width', 'height', 'color_score']
x=df[feature_names]
y=df['fruit_label']
x_train, x_test, y_train, y_test = train_test_split(x,y, random_state=0)
print(x_train[:3]) # to check output
scaler = MinMaxScaler()
x_train=scaler.fit_transform(x_train)
x_test= scaler.transform(x_test)
print("\nAfter scaling\n")
print(x_train[:3])

```

	mass	width	height	color_score
42	154	7.2	7.2	0.82
48	174	7.3	10.1	0.72
7	76	5.8	4.0	0.81

After scaling

```

[[0.27857143 0.41176471 0.49230769 0.72972973]
 [0.35      0.44117647 0.93846154 0.45945946]
 [0.       0.       0.       0.7027027 ]]

```

```
from sklearn.linear_model import LogisticRegression # machine learning  
lib/model
```

```
feature_names = ['mass', 'width', 'height', 'color_score']
```

```
x=df[feature_names]
```

```
y=df['fruit_label']
```

```
x_train, x_test, y_train, y_test = train_test_split(x,y, random_state=0)
```

```
scaler = MinMaxScaler()
```

```
x_train=scaler.fit_transform(x_train)
```

```
x_test= scaler.transform(x_test)
```

```
#logistic regression
```

```
logreg = LogisticRegression() # machine learning algorithm
```

```
logreg.fit(x_train, y_train)
```

```
#print score of train data
```

```
print('Accuracy of Logistic regression classifier on training set:{:.2f}'
```

```
      .format(logreg.score(x_train, y_train)))
```

```
#print score of test data
```

```
print('Accuracy of Logistic regression classifier on test set:{:.2f}'
```

```
      .format(logreg.score(x_test, y_test)))
```

```
Accuracy of Logistic regression classifier on training set:0.75
```

```
Accuracy of Logistic regression classifier on test set:0.47
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# KNN method

knn = KNeighborsClassifier()

knn.fit(x_train, y_train)


# print score of train data
print('Accuracy of KNN classifier on training set: {:.2f}'
      .format(knn.score(x_train, y_train)))


# print score of test data
print('Accuracy of KNN Classifier on test set: {:.2f}'
      .format(knn.score(x_test, y_test)))

Accuracy of KNN classifier on training set:0.95
Accuracy of KNN Classifier on test set:1.00

from sklearn.svm import SVC


# SVM classifier

svm = SVC()

svm.fit(x_train, y_train)


# print score of train data
print('Accuracy of SVM classifier on training set: {:.2f}'
      .format(svm.score(x_train, y_train)))


# print score of test data
```

```
print('Accuracy of SVM Classifier on test set:{:.2f}'
```

```
      .format(svm.score(x_test, y_test)))
```

Accuracy of SVM classifier on training set:0.91

Accuracy of SVM Classifier on test set:0.80

```
data = {'Training Accuracy (in %)':[75,95,91],'Testing Accuracy (in %)':[47,100,80]}
```

```
df1 = pd.DataFrame(data, index =['Logistic Regression','K-Nearest Neighbour (KNN)','Support Vector Machine (SVM)'])
```

```
df1
```

	Training Accuracy (in %)	Testing Accuracy (in %)
Logistic Regression	75	47
K-Nearest Neighbour (KNN)	95	100
Support Vector Machine (SVM)	91	80


## Practical No – 2

2) On the iris dataset, perform KNN algorithm and discuss result

```
import pandas as pd
iris = pd.read_csv("iris.csv")
iris.dtypes

import matplotlib.pyplot as plt # mostly used for visualization purposes
import numpy as np
import seaborn as sns

sns.distplot(iris['Sepal.Length'], kde=False,label='Sepal.Length')
sns.distplot(iris['Sepal.Width'], kde=False,label='Sepal.Width')
sns.distplot(iris['Petal.Length'], kde=False,label='Petal.Length')
sns.distplot(iris['Petal.Width'], kde=False,label='Petal.Width')
plt.legend()
```



```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

```
feature_names = ['Sepal.Length', 'Sepal.Width', 'Petal.Length', 'Petal.Width']
x=iris[feature_names]
y=iris['Species']
```

```
x_train, x_test, y_train, y_test = train_test_split(x,y, random_state=0)
```

```
print(x_train[:3]) # to check output
```

```
scaler = MinMaxScaler()
x_train=scaler.fit_transform(x_train)
x_test= scaler.transform(x_test)
```

```
print("\nAfter scaling\n")
```

```
print(x_train[:3]) # to check output
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
61	5.9	3.0	4.2	1.5
92	5.8	2.6	4.0	1.2
112	6.8	3.0	5.5	2.1

After scaling

```
[[0.44444444 0.41666667 0.53448276 0.58333333]
```

```
[0.41666667 0.25    0.5    0.45833333]
```

```
[0.69444444 0.41666667 0.75862069 0.83333333]]
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# KNN method
```

```
knn = KNeighborsClassifier()
```

```
knn.fit(x_train, y_train)
```

```
#print score of train data
```

```
print('Accuracy of KNN classifier on training set:{:.2f}'  
      .format(knn.score(x_train, y_train)))
```

```
#print score of test data
```

```
print('Accuracy of KNN Classifier on test set:{:.2f}'  
      .format(knn.score(x_test, y_test)))
```

```
Accuracy of KNN classifier on training set:0.96
```

```
Accuracy of KNN Classifier on test set:0.97
```



### Practical No – 3

3) Implement apriori algorithm on online retail dataset and discuss result

```
import os

for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

/kaggle/input/online-retail-ii-uci/online_retail_II.csv

transaction_df= pd.read_csv("../input/online-retail-ii-uci/online_retail_II.csv")
transaction_df

transaction_df = transaction_df[transaction_df.Country=='France']
transaction_filtered = transaction_df[['Invoice','Description','Quantity']].copy()
transaction_filtered

transaction_filtered.sort_values(by='Quantity', ascending=True)
transaction_filtered = transaction_filtered[transaction_filtered.Quantity > 0 ]
transaction_filtered.sort_values(by='Quantity', ascending=True)
transaction_filtered['Quantity']= [1]*len(transaction_filtered)
invoice = list(transaction_filtered.Invoice)

index_no = [invoice[index] for index in np.arange(len(invoice)) if not
invoice[index].isnumeric()]

transaction_filtered[transaction_filtered['Invoice'].isin(index_no)]

transaction_filtered=
transaction_filtered[~transaction_filtered['Invoice'].isin(index_no)]

invoice = list(transaction_filtered.Invoice)

index_no = [index for index in np.arange(len(invoice)) if not
invoice[index].isnumeric()]

transaction_filtered.iloc[index_no,:]
```

```

temp_df = transaction_filtered[transaction_filtered.Description !=
transaction_filtered.Description]

temp_df

for invoice in list(temp_df.Invoice):

    if len(transaction_filtered[transaction_filtered.Invoice == invoice]) > 1:

        print((str)(invoice))

        temp = transaction_filtered[transaction_filtered.Invoice ==
invoice].groupby(['Invoice']).agg({'Description':lambda x: list(x)})

        if len(list(set(temp)))>0 :

            print(temp)

transaction_filtered.dropna(axis=0, inplace=True)

transaction_filtered

def return_one(x):

    return 1

table = pd.pivot_table(transaction_filtered, values='Quantity', index=['Invoice'],
                        columns=['Description'], aggfunc=return_one, fill_value=0)

table

frequent_itemsets = apriori(table, min_support=0.01, use_colnames=True)

frequent_itemsets

rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)

rules

rules.sort_values(by=['support','confidence'], ascending=False)

```

## Practical No – 4

4)Implement Naïve Bayes Classifier and K-Nearest Neighbor Classifier on Data set of your choice. Test and Compare for Accuracy and Precision.

```
df=pd.read_csv('fruit_data.csv')
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

feature_names = ['mass', 'width', 'height', 'color_score']
x=df[feature_names]
y=df['fruit_label']

x_train, x_test, y_train, y_test = train_test_split(x,y, random_state=0)

print(x_train[:3]) # to check output

scaler = MinMaxScaler()
x_train=scaler.fit_transform(x_train)
x_test= scaler.transform(x_test)

print("\nAfter scaling\n")
print(x_train[:3]) # to check output
    mass width height color_score
42  154   7.2   7.2     0.82
48  174   7.3  10.1     0.72
```

7    76   5.8   4.0       0.81

After scaling

```
[[0.27857143 0.41176471 0.49230769 0.72972973]
```

```
[0.35       0.44117647 0.93846154 0.45945946]
```

```
[0.       0.       0.       0.7027027 ]]
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# KNN method
```

```
knn = KNeighborsClassifier()
```

```
knn.fit(x_train, y_train)
```

```
#print score of train data
```

```
print('Accuracy of KNN classifier on training set:{:.2f}'
```

```
      .format(knn.score(x_train, y_train)))
```

```
#print score of test data
```

```
print('Accuracy of KNN Classifier on test set:{:.2f}'
```

```
      .format(knn.score(x_test, y_test)))
```

```
Accuracy of KNN classifier on training set:0.95
```

```
Accuracy of KNN Classifier on test set:1.00
```

```
from sklearn.naive_bayes import GaussianNB
```

```
# Gaussian Naive bayes
```

```
gnb = GaussianNB()
```

```
gnb.fit(x_train, y_train)
```

```
#print score of train data
```

```
print('Accuracy of GNB classifier on training set:{:.2f}'
```

```
.format(gnb.score(x_train, y_train)))
```

```
#print score of test data
```

```
print('Accuracy of GNB Classifier on test set:{:.2f}'
```

```
.format(gnb.score(x_test, y_test)))
```

```
Accuracy of GNB classifier on training set:0.86
```

```
Accuracy of GNB Classifier on test set:0.67
```

```
apply confusin matrix
```

```
confusion matrix for KNN
```

```
import sklearn
```

```
from sklearn.metrics import confusion_matrix
```

```
import seaborn as sns
```

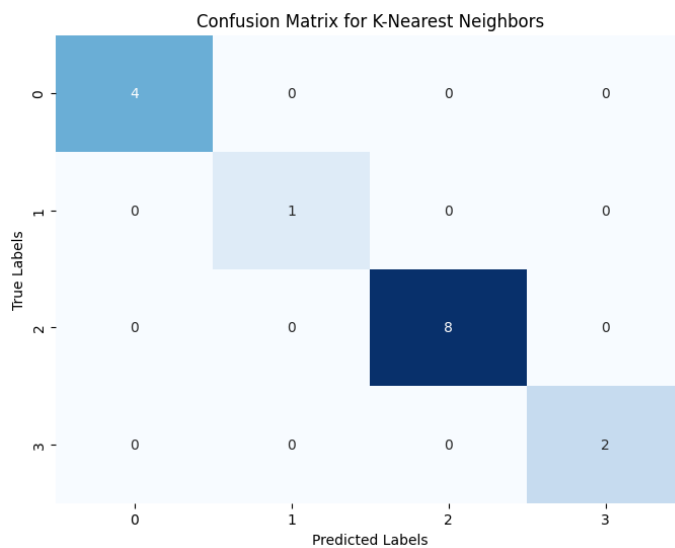
```
import matplotlib.pyplot as plt
```

```
# Assuming you have already trained your classification models, e.g., K-Nearest  
Neighbors (knn)
```

```
y_pred_knn = knn.predict(x_test) # Make predictions on the test data
```

```
# Create the confusion matrix for K-Nearest Neighbors
confusion_knn = confusion_matrix(y_test, y_pred_knn)

# Create a heatmap of the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_knn, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for K-Nearest Neighbors')
plt.show()
```



```
confusion matrix for naive bayes

from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming you have already trained your classification models
```

```
y_pred_gnb = gnb.predict(x_test) # Make predictions on the test data
```

```
# Create the confusion matrix
```

```
confusion_gnb = confusion_matrix(y_test, y_pred_gnb)
```

```
# Create a heatmap of the confusion matrix
```

```
plt.figure(figsize=(8, 6))
```

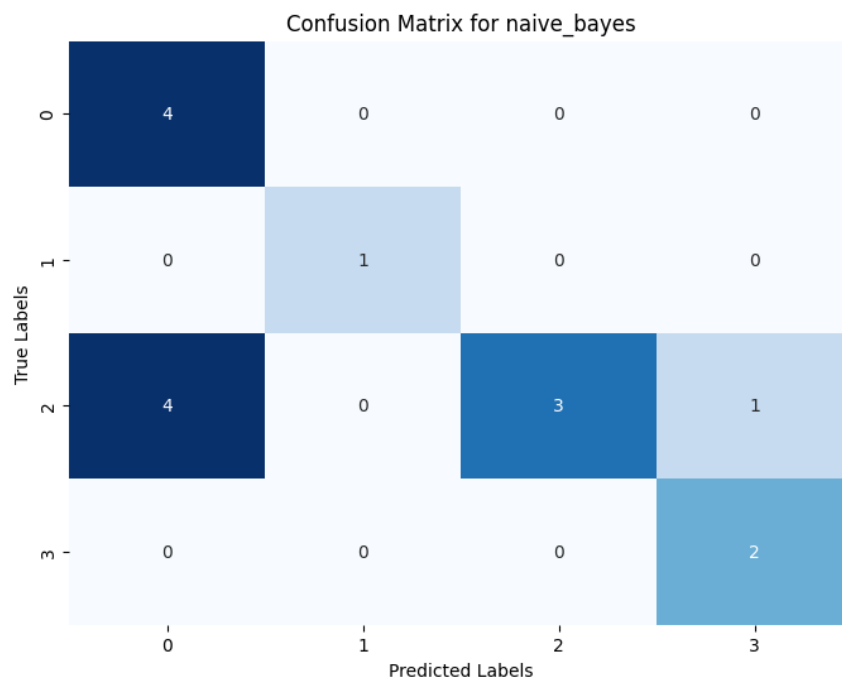
```
sns.heatmap(confusion_gnb, annot=True, fmt='d', cmap='Blues', cbar=False)
```

```
plt.xlabel('Predicted Labels')
```

```
plt.ylabel('True Labels')
```

```
plt.title('Confusion Matrix for naive_bayes')
```

```
plt.show()
```




## Practical No – 5

5) Implement K-means Clustering on a proper dataset of your choice


**K-Means Clustering :** Perform clustering for the crime data and identify the number of clusters formed and draw inferences. Refer to crime\_data.csv dataset.

```
crime = pd.read_csv("crime_data.csv")
crime['State'] = crime.iloc[:,0]
crime = crime.iloc[:, [5,1,2,3,4]]
crime.isna().sum()
crime1 = crime.duplicated()
sum(crime1)
plt.boxplot(crime.Murder);plt.title('Boxplot');plt.show()
```




```
plt.boxplot(crime.Assault);plt.title('Boxplot');plt.show()
```






```
plt.boxplot(crime.Rape);plt.title('Boxplot');plt.show()
```



```
from scipy.stats.mstats import winsorize  
crime['Rape'] = winsorize(crime.Rape, limits=[0.07, 0.093])  
plt.boxplot(crime['Rape']);plt.title('Boxplot');plt.show()
```



```
(crime == 0).all()
```

```
State      False
```

```
Murder     False
```

```
Assault    False
```

```
UrbanPop   False
```

```
Rape       False
```

```
dtype: bool
```

```
from sklearn.cluster import KMeans
```

```
TWSS = []
```

```
k = list(range(2, 8))
```

```
for i in k:
```

```
    kmeans = KMeans(n_clusters = i)
```


```
    kmeans.fit(df_norm)
```

```
    TWSS.append(kmeans.inertia_)
```

TWSS

```
[7.358376498536079,  
 5.532071995078602,  
 4.0407678952238815,  
 3.5539811127025747,  
 3.1628651131109455,  
 2.8417637970747243]
```

```
plt.plot(k, TWSS, 'ro-');plt.xlabel("No_of_Clusters");plt.ylabel("total_within_SS")
```



```
model = KMeans(n_clusters = 4)  
model.fit(df_norm)  
KMeans(n_clusters=4)  
crime.iloc[:, 1:6].groupby(crime.clust).mean()
```